

Coming from superpowers?

You already know 80% of chris-code.

Same `brainstorm → plan → execute → review → finish` pipeline.
Same skill names. Same TDD and systematic-debugging discipline.

This deck is the other **20%**, and the *why*.

What feels different, day to day

1

Shorter specs & plans

No more 2,500-word specs or 10k-word plans full of code the implementer throws away.

2

Dedicated agents

Coding and review run through named agents auto-dispatched by file type. You rarely pick one.

3

Same gates, every task

Spec compliance → code quality → pre-commit lint. Nothing is "too small to review."

Six thematic shifts

- 1 **Lean artifacts** over exhaustive ones
- 2 **A real agent layer**, dispatched by scope
- 3 **Review** is a uniform, multi-stage gate
- 4 **Parallelism** is a feature, not a footgun
- 5 **Native tools** first, with a hard gate
- 6 **Coherence enforcement** — every determined change fits the codebase (*the largest divergence*)

1 • Lean artifacts over exhaustive ones

superpowers

Spec written "as if zero context"

Plan: complete code in every step

specs grew past 2,500 w • plans 5,000–10,000 w

agents discarded the pasted code and rewrote it



chris-code

Spec = invariants only


Plan = what & where, no code

code written against the real repo

"contracts stay, choreography goes"

2 • A real agent layer, dispatched by scope

superpowers dispatches a generic `Task (general-purpose)` subagent steered by a prompt template. chris-code routes to dedicated agents:



3 • Review: kept the spine, hardened it



Same two-stage spine. chris-code makes the reviewers **dedicated, scope-dispatched agents**, adds a **mandatory lint gate** and a **spec-blind intent re-check**, and re-reads the actual code at every gate. More gates raise *recall*, not proof — the decorrelated axes (deterministic lint, spec-blind intent, mutation probe) carry the weight, not repetition.

4 • Parallelism is a feature, not a footgun

✗ superpowers — "dispatch parallel implementers" is a Red Flag

implementer A

implementer B

implementer C

unscheduled: they edit the same files and collide

✓ chris-code — map file footprints, then stage

Stage 1

A • api/

B • db/

C • ui/

Stage 2

D • api/ + db/ (waits — overlaps A & B)

5 • Native tools first, with a hard gate



The split is only at "no native tool": superpowers falls back; chris-code refuses, because a silent fallback once caused real damage.

6 • Coherence enforcement – the largest divergence

superpowers

A review refactors in place

A fix is patched directly

A determined change is just coded

the first impl that works lands – even when the repo already has a pattern



chris-code

One **application engine**: `coherent-change`

Every determined change routes through it
defend → size → (batch a whole set)

discovery proposes & hands off – it never lands its own patches

superpowers has no step that makes a change **fit the codebase it lands in**. chris-code adds one: discover candidates from the existing code, defend the most coherent against the alternatives (a correctness table + a real rebuttal of each), **size** the work (build inline, or route to spec → plan → execution), and **batch** a whole set of findings into one coordinated change.

What stayed the same

- ✓ **brainstorming** — intent → approaches → design
- ✓ **test-driven-development** — RED-GREEN-REFACTOR
- ✓ **systematic-debugging** — four phases, Iron Law
- ✓ **finishing-a-development-branch** — merge / PR / keep / discard
- ✓ **dispatching-parallel-agents, receiving-code-review, using-chris-code** (light edits)


Skills you know – what changed most

Where muscle memory will mislead you. Framed before → after:

Skill	superpowers	chris-code
writing-plans	the plan skill: exhaustive, full code per step	slimmed to <code>lean-plan</code> ; spec → <code>lean-spec</code>
subagent-driven-development	2-stage review; parallel discouraged	3 gates/task, scope dispatch, staged parallelism
verification-before-completion	"what command proves this? run it"	6 steps: tests → lints → review → requirements → spec-blind intent re-check → mutation re-check
requesting-code-review	primary, mandatory path	demoted to ad-hoc; base <code>HEAD~1</code> → <code>merge-base HEAD main</code>

What's new, and why

New skills (11)



New agents (14) – the layer superpowers doesn't have

3 coders (`python` / `pytorch` / `rust`) – exclusive, most-specific wins

3 quality-reviewers – additive post-spec review

2 review-lite gates – pre-commit idiom + lint

2 design-reviewers (`python` / `rust`) – senior cohesion at the verification gate

2 conformance reviewers – `spec-reviewer` (promoted from a prompt template) + `intent-reviewer` (spec-blind, new)

`bug-hunter` – adversarial edge-case test writer

`mutation-tester` – polyglot mutation gate in an isolated worktree (Step 6 / on-demand)

The net

chris-code is a **true superset**: every superpowers skill is present (one renamed, one split), plus 11 new skills and a 14-agent layer.

Same spine. Leaner artifacts, mechanical dispatch, harder gates.

Full write-up: [superpowers-comparison.md](#)