

chris-code

An opinionated Claude Code workflow

Design before it codes. Review before it lands.

25 skills · 13 agents · uniform review gates

What it is

A personal **Claude Code plugin** that turns the assistant from a free-form chat tool into a **fixed engineering pipeline**.

Drop it into any repo. Describe a feature (or type `/brainstorming`) and the workflow takes over:

“ brainstorm intent → write a lean spec → hand off a thin plan → dispatch a coder per task → review in stages → gate every commit. ”

Derived from [obra/superpowers](https://github.com/obra/superpowers), redesigned around four failures it kept hitting:

Why it exists: four failures

1

Spec & plan bloat

Specs grew past 2,500 words; plans piled on another 5,000–10,000 words of code.

2

Shadow code

Plans mandated complete code per step. Agents discarded it and rewrote against the real repo.

3

Inconsistent review

Review folded into a single self-check pass, applied unevenly across tasks.

4

No coherence enforcement

Changes patched in place. The first impl that works lands — even when the repo already has a pattern for it.

“ The rest of this deck is how chris-code answers each. The fourth — the largest divergence — is the engine on the next-to-last slide. ”

The core idea

"Contracts stay, choreography goes."

Contracts stay

Spec invariants

Public APIs & signatures

Data schemas

spec: contracts only

Choreography goes

~~Step-by-step code~~


~~Inline scaffolds~~

~~"do this, then that"~~

plan: what/where, no code

Agents ignored prescriptive code blocks and rewrote them anyway. Plans now say **what** and **where**; code is written against the real repo.

The pipeline



The agent layer

a .py file in a torch project

↓ dispatched by scope (extension + deps) ↓

coder — exclusive

pytorch-coder ✓

python-coder

most-specific wins

quality-reviewer — additive

python-quality-reviewer

pytorch-quality-reviewer

all matching fire

review-lite — commit gate

python-review-lite


idiom + lint, per commit

design-reviewer — verification gate

python-design-reviewer

senior cohesion, read-only

Review is a uniform, multi-stage gate



The **same** gates on every task. No task is "small enough to skip." **Do Not Trust the Report** — reviewers re-read the code, not the summary. More stages raise *recall*, not proof: they catch more, they don't certify the rest is clean.

Parallelism is a feature, not a footgun

Stage 1

Task A • api/

Task B • db/

Task C • ui/


file footprints mapped up front • non-overlapping tasks run concurrently

Stage 2

Task D • api/ + db/ (waits — overlaps A & B)

Superpowers flagged parallel implementers as dangerous. chris-code solves it by **scheduling**, not prohibition.

Verification before "done"



No completion claim without fresh evidence. "I'm confident" → run the commands. **Green means these lenses caught nothing, not that nothing's wrong** — the independent axes (deterministic lint, spec-blind intent) carry more than another same-model re-read.


Beyond the pipeline: quality campaigns

On-demand skills for when you want to go hunting:

- `bug-hunt` — parallel adversarial edge-case tests across subsystems
- `test-sweep` — iterative combinatorial test-and-fix campaign
- `code-archaeology` — dead code, stubs, spec-vs-impl gaps
- `technical-review` — math / algorithm / numerical correctness for ML
- `python-review` • `rust-review` — senior refactor & API-design passes
- `release` — version bump + changelog + GitHub release

One engine for every determined change

A bug, a refactor, an API alignment, a finding from a review — the behavior is settled, only the **implementation** is open. chris-code routes them all through a single **application engine**.



Discovery is separated from application: audits and reviews propose end-states and hand off — they never land their own patches. Signature output: a **defended choice** — every candidate rooted in how the codebase already solves it, plus why the chosen one wins.

Size the change – or batch the set

After the defended choice, the engine **sizes the work** and routes it – so no determined change is built bigger than it should be, or smaller.

Single coherent edit

Build **inline** – coder agent →

`*-review-lite` self-gate → hand back.

The caller runs the close.

Major / multi-task

Settled design, not a one-shot build –
route it to planned execution:

`lean-spec` → `lean-plan` →
`subagent-driven-development` .

A set of findings – batch

One consolidated research pass, a
defended choice per change, the whole
set into **one** spec → plan → SDD.

Same defend-and-prove discipline at every scale. A change **closes its whole scope** – every sibling case its intent reaches – and defers only what's genuinely separate. Ship the change that looks like it was always there.

How to use it

Drop into any repo. Describe a feature, or type `/brainstorming`.

The workflow designs before it codes, keeps the main context clean by offloading to focused subagents, and **catches agent drift before it reaches** `main`.

That's chris-code.